

```
In [1]: import torch
import torch.nn as nn
```

0. Prepare Input.

```
In [2]: audio = torch.rand(1, 160000) # 10s audio over 16kHz SR.
x = audio[None].repeat(2, 1, 1) # make it a batch.
x.shape
```

```
Out[2]: torch.Size([2, 1, 160000])
```

1. Feature Extractor (Downsample with Conv).

```
In [3]: def conv_block(down_rate, c_in=512, c_out=512, kernel_size=3, **kwargs):
    assert kernel_size % 2, 'For simplicity, make sure kernel_size is odd.'
    return nn.Sequential(
        nn.Conv1d(
            in_channels=c_in,
            out_channels=c_out,
            stride=down_rate,
            kernel_size=kernel_size,
            padding=kernel_size // 2,
            **kwargs,
        ),
        nn.GELU()
    )

feature_extractor = torch.nn.Sequential(
    conv_block(5, 1, 512, kernel_size=9),
    *[conv_block(down_rate) for down_rate in [2,] * 6]
)

x = feature_extractor(x)
x.shape
```

```
Out[3]: torch.Size([2, 512, 500])
```

2.1 Vector-Quantization

Why VQ? Enforce continuous vector to be some limited representations (Codebook) and therefore prevent overfitting.

What is a Codebook? A trainable group of vectors.

```
In [4]: from einops import rearrange, einsum

class VectorQuantizer(nn.Module):
    def __init__(self, n_group=2, group_len=320, n_dim=128, feat_in=512):
        super().__init__()
        self.n_group = n_group
        self.scorer = nn.Linear(feat_in, n_group * group_len)
        self.codebook = nn.Parameter(torch.randn(n_group, group_len, n_dim))

    def forward(self, x):
        x = self.scorer(x) # 'b l (n_group group_len)'
        score = rearrange(x, 'b l (g n) -> b l g n', g=self.n_group)
        score = score.softmax(dim=-1)[..., None] # (b l g n 1)

        codebook = self.codebook[None, None] # (1 1 g n d)

        q = rearrange((score * codebook).sum(dim=-2), 'b l g d -> b l (g d)')
        return q
```

```
In [5]: _out = VectorQuantizer()(rearrange(x, 'b d l -> b l d'))
        _out.shape
```

```
Out[5]: torch.Size([2, 500, 256])
```

Why Gumbel?

Probability Sampling. Here is a example.

```
In [6]: p = torch.tensor([0.6, 0.3, 0.1])
        for _ in range(10):
            print(p.argmax().item(), end=' ')
```

0 0 0 0 0 0 0 0 0 0

```
In [7]: def gumbel(p):
        gumbel_noise = - torch.log(- torch.log(torch.rand_like(p)))
        return p + gumbel_noise

        for _ in range(10):
            print(gumbel(p).argmax().item(), end=' ')
```

0 2 0 1 0 2 0 0 1 2

Gumbel Vector-Quantier

```
In [8]: class GumbelVectorQuantizer(nn.Module):
        def __init__(self, n_group=2, group_len=320, n_dim=128, feat_in=512):
            super().__init__()
            self.n_group = n_group
            self.scorer = nn.Linear(feat_in, n_group * group_len)
            self.codebook = nn.Parameter(torch.randn(n_group, group_len, n_dim))

        def gumbel_softmax(self, p, dim, tau=0.1, eps=1e-8):
            gumbel_noise = - torch.log(- torch.log(torch.rand_like(p) + eps) + eps)
            return ((p + gumbel_noise) / tau).softmax(dim=dim)

        def forward(self, x):
            x = self.scorer(x)  # 'b l (n_group group_len)'
            score = rearrange(x, 'b l (g n) -> b l g n', g=self.n_group)
            score = self.gumbel_softmax(score, dim=-1)[..., None]  # (b l g n 1)

            codebook = self.codebook[None, None]  # (1 1 g n d)

            q = rearrange((score * codebook).sum(dim=-2), 'b l g d -> b l (g d)')
            return q
```

```
In [9]: gumbel_vq = GumbelVectorQuantizer()
        q = gumbel_vq(rearrange(x, 'b d l -> b l d'))
```

```
q.shape
```

```
Out[9]: torch.Size([2, 500, 256])
```

2.2 Mask Feature.

Mask Feature at time dimension.

```
In [10]: class FeatureMasker(nn.Module):
    def __init__(self, feature_dim=512, n_masks=8, mask_len=10):
        super().__init__()
        # replacer is a learnable vector, not zero vector.
        self.vec_replacer = nn.Parameter(torch.randn(feature_dim))
        self.n_masks = n_masks
        self.mask_len = mask_len

    def random_mask(self, x, fill_value=None):
        b = x.shape[0]
        start_points = torch.randint(0, x.shape[1] - self.mask_len, (b, self.n_masks,))
        end_points = start_points + self.mask_len
        ref = torch.zeros(b, x.shape[1])
        for i in range(self.n_masks):
            s, e = start_points[:, i], end_points[:, i]
            for b_i in range(b):
                ref[b_i, s[b_i]: e[b_i]] = 1

        idx = (ref > 0).nonzero()
        x[idx[:, 0], idx[:, 1]] = self.vec_replacer

        return x, idx

    def forward(self, x):
        # x -> (b l d)
        return self.random_mask(x)
```

```
In [11]: masker = FeatureMasker()
x, mask_idx = masker(x.permute(0, 2, 1))
x.shape, mask_idx.shape
```

```
Out[11]: (torch.Size([2, 500, 512]), torch.Size([157, 2]))
```

2.3 Transformer Encoder (Model Long-range Correlation)

```
In [12]: class MHAttn(nn.Module):
    def __init__(self, dim=768, n_heads=12):
        super().__init__()
        self.n_heads = n_heads
        self.to_qkv = nn.Linear(dim, dim * 3)
        self.to_out = nn.Linear(dim, dim)
        self.d_root = dim ** 0.5
        self.to_mh = lambda x: rearrange(x, 'b l (h d) -> (b h) l d', h=self.n_heads)
        self.mh_to_d = lambda x: rearrange(x, '(b h) l d -> b l (h d)', h=self.n_heads)

    def forward(self, x):
        q, k, v = list(map(self.to_qkv, x).chunk(3, dim=-1))
        attn = (einsum(q, k, 'B i d, B j d -> B i j') / self.d_root).softmax(dim=1)
        return self.to_out(self.mh_to_d(attn @ v))

class FF(nn.Module):
    def __init__(self, dim=768):
        super().__init__()
        self.ln = nn.LayerNorm(normalized_shape=dim)
        self.up = nn.Sequential(
            nn.Linear(dim, dim * 4),
            nn.GELU()
        )
        self.down = nn.Sequential(
            nn.Linear(dim * 4, dim),
            nn.GELU()
        )

    def forward(self, x):
        return self.down(self.up(self.ln(x)))

class TransformerBlock(nn.Module):
    def __init__(self, dim=768):
        super().__init__()
```

```

        self.attn = MHAttn(dim)
        self.ff = FF(dim)

    def forward(self, x):
        return x + self.ff(self.attn(x))

```

```

In [13]: to_transformer_dim = nn.Conv1d(512, 768, kernel_size=1)
        pe = conv_block(1, c_in=768, c_out=768, kernel_size=127, groups=12)
        transformer_encoder = nn.Sequential(*[TransformerBlock() for _ in range(12)])

```

```

In [14]: x = rearrange(x, 'b d l -> b l d')
        c = to_transformer_dim(x)
        c = c + pe(c)
        c = rearrange(c, 'b d l -> b l d')
        c = transformer_encoder(c)
        c.shape

```

```

Out[14]: torch.Size([2, 500, 768])

```

3. Self-Supervised Learning

```

In [15]: import torch.nn.functional as F

```

```

In [16]: to_vq_dim = nn.Linear(768, 256)
        c = to_vq_dim(c)

```

```

In [17]: c_pred_masked = F.normalize(c[mask_idx[:,0], mask_idx[:, 1]], dim=-1)
        q_masked = F.normalize(q[mask_idx[:,0], mask_idx[:, 1]], dim=-1)
        c_pred_masked.shape, q_masked.shape

```

```

Out[17]: (torch.Size([157, 256]), torch.Size([157, 256]))

```

3.1 Main Loss: Predict (from Transformer Encoder) v.s. Qutanzed Vecotr

```

In [18]: # Similarity Matrix: [N_mask, N_mask]
        logits = c_pred_masked @ q_masked.T
        labels = torch.arange(logits.shape[0], device=logits.device) # only pred in diognal is true.

```

```
temperature = 1.0

loss_nce = torch.nn.functional.cross_entropy(logits / temperature, labels)
```

3.2 Extra Loss: Codebook Usage; Feature Penalty; ...More

3.2.1 Codebook Usage

In [19]: *# Encourage Codebook Usage*

```
class GumbelVectorQuantizer_(GumbelVectorQuantizer):
    """re-implement it to obtain score (of each vec in the code book)"""
    def forward(self, x):
        x = self.scorer(x) # 'b l (n_group group_len))
        score = rearrange(x, 'b l (g n) -> b l g n', g=self.n_group)
        score = self.gumbel_softmax(score, dim=-1)[..., None] # (b l g n 1)

        codebook = self.codebook[None, None] # (1 1 g n d)
        q = rearrange((score * codebook).sum(dim=-2), 'b l g d -> b l (g d)')

        return q, score
```

In [20]: `vq_ = GumbelVectorQuantizer_()`
`q_, score = vq_(rearrange(x, 'b d l -> b l d'))`

In [21]: `prob_per_vec = score[..., 0].mean(dim=(0, 1, 2))`
`prob_per_vec.shape`

Out[21]: `torch.Size([320])`

In [22]: `loss_codebook = (prob_per_vec * torch.log(prob_per_vec + 1e-9)).sum()`
`loss_codebook`

Out[22]: `tensor(-5.6982, grad_fn=<SumBackward0>)`

3.2.3 Feature Penalty

```
In [23]: # `x` is obtained from (Conv) `feature_extractor(wave)`  
loss_feature_pen = (x ** 2).mean()  
loss_feature_pen
```

```
Out[23]: tensor(0.1466, grad_fn=<MeanBackward0>)
```

3.3 Loss for Training.

```
In [24]: total_loss = loss_nce + 0.1 * loss_codebook + 10 * loss_feature_pen  
total_loss
```

```
Out[24]: tensor(5.9539, grad_fn=<AddBackward0>)
```

4. Finally!

After self-supervised training, we can use trained CNN + Transformer-Encoder as a powerful feature extractor for downstream tasks.